


CS826




Neural Networks

Week 1 - Summary




Week 1: Review and Getting Started with Neural Networks

This week you will be introduced to:


1. What is deep learning? 
2. Motivation, challenges, and applications 
3. Architecture and maths 
4. Simple examples
5. Google Colab for practical sessions

How do ANNs work?

Biological Neuron



Artificial Neuron



An artificial neuron is an imitation of a human neuron

Forward Propagation - Perceptron

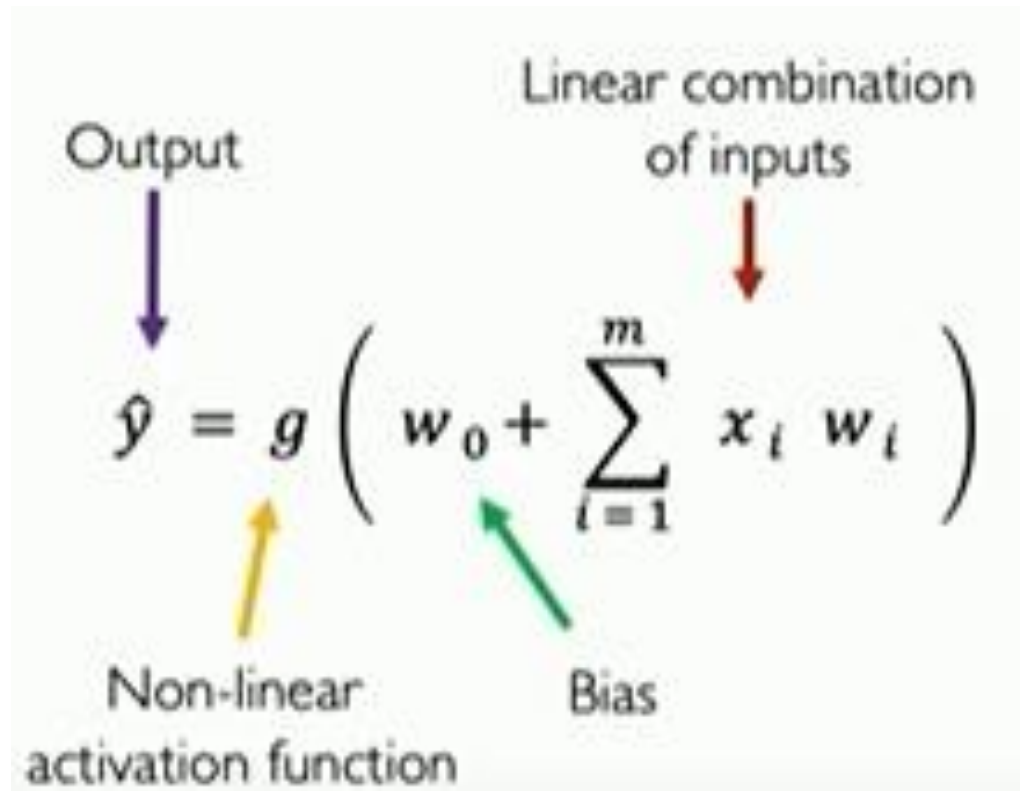
Output

Linear combination of inputs


$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function

Bias



The diagram shows the mathematical formula for forward propagation in a perceptron. The output \hat{y} is calculated by applying a non-linear activation function g to a linear combination of inputs. The linear combination consists of a bias term w_0 and the weighted sum of input features x_i with weights w_i . Annotations include a purple arrow pointing to \hat{y} labeled 'Output', a red arrow pointing to the sum term labeled 'Linear combination of inputs', a yellow arrow pointing to g labeled 'Non-linear activation function', and a green arrow pointing to w_0 labeled 'Bias'.

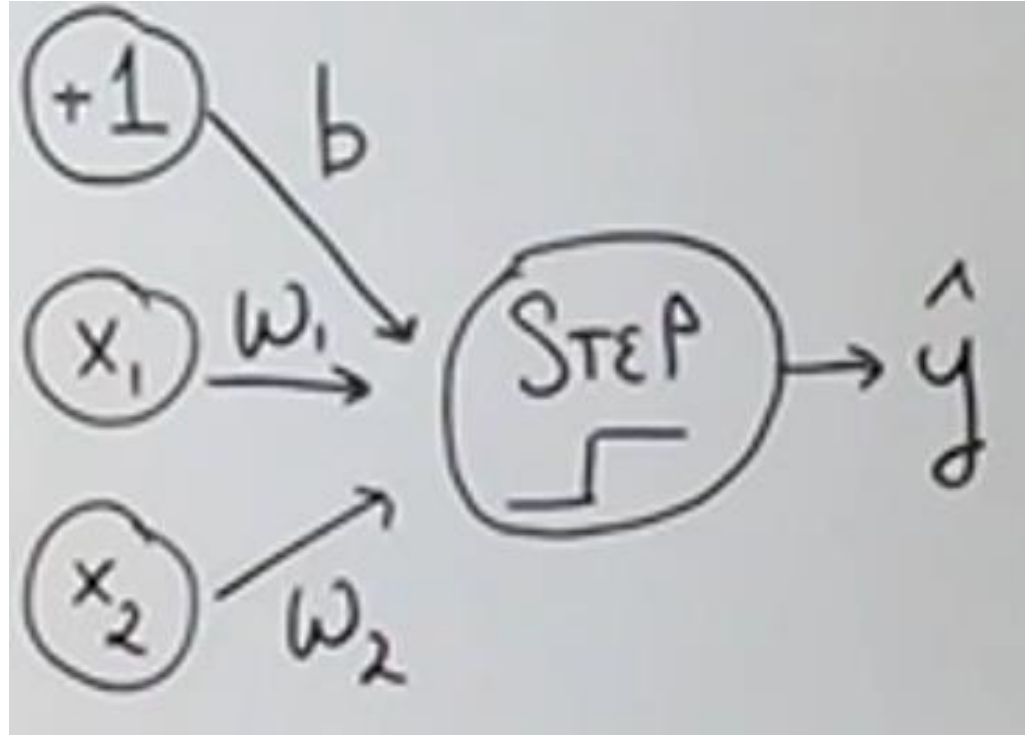


Perceptron

Let the input feature vector be:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad y \in \{1, 0\}$$

where y is the binary response.



The perceptron model computes a weighted sum of inputs:

$$z = w_1 x_1 + w_2 x_2 + b$$

The perceptron then applies a step function to determine the predicted label \hat{y} :


$$\hat{y} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

Logistic Regression

Let the input feature vector be:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad y \in \{1, 0\}$$

where y is the binary response.



In logistic regression, we compute the same linear combination of inputs:

$$z = w_1x_1 + w_2x_2 + b$$

but instead of a step function, we apply the sigmoid function $\sigma(z)$ to get the predicted probability \hat{p} :

$$\hat{p} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

Neural Network


$$z_1^{(1)} = W_{11}X_1 + W_{21}X_2 + b_1^{(1)}$$

$$z_2^{(1)} = W_{12}X_1 + W_{22}X_2 + b_2^{(1)}$$

$$h_1 = \sigma(z_1^{(1)}) \quad \text{and} \quad h_2 = \sigma(z_2^{(1)})$$

$$z^{(2)} = w_1^{(2)}h_1 + w_2^{(2)}h_2 + b^{(2)}$$

$$\hat{p} = \sigma(z^{(2)})$$



Matrix Form: $\mathbf{z}^{(1)} = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$

$$\mathbf{W}^{(1)} = \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix}, \quad \mathbf{b}^{(1)} = \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \end{bmatrix}$$

$$\mathbf{h} = \sigma(\mathbf{z}^{(1)})$$

$$\mathbf{z}^{(2)} = \mathbf{w}^{(2)}\mathbf{h} + b^{(2)}$$


$$\hat{p} = \sigma(z^{(2)})$$

Vital and Lingering Questions

Some questions regarding neural network architecture and training include:


- How to get weights and biases (w and b)?
- How many hidden layers should be used?
- What activation function to choose?

Activation Functions



Importance of Activation Functions

Distinguish Red vs Green points



Introduce
non-linearities
to the network

What is bias, how is used
and why it is important?

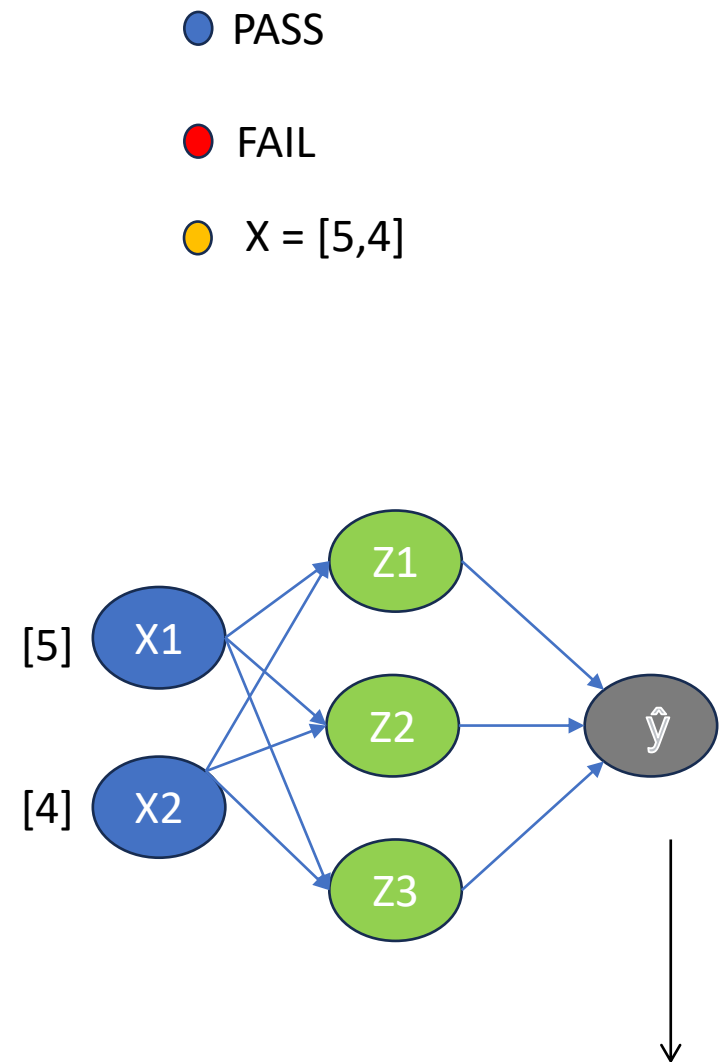
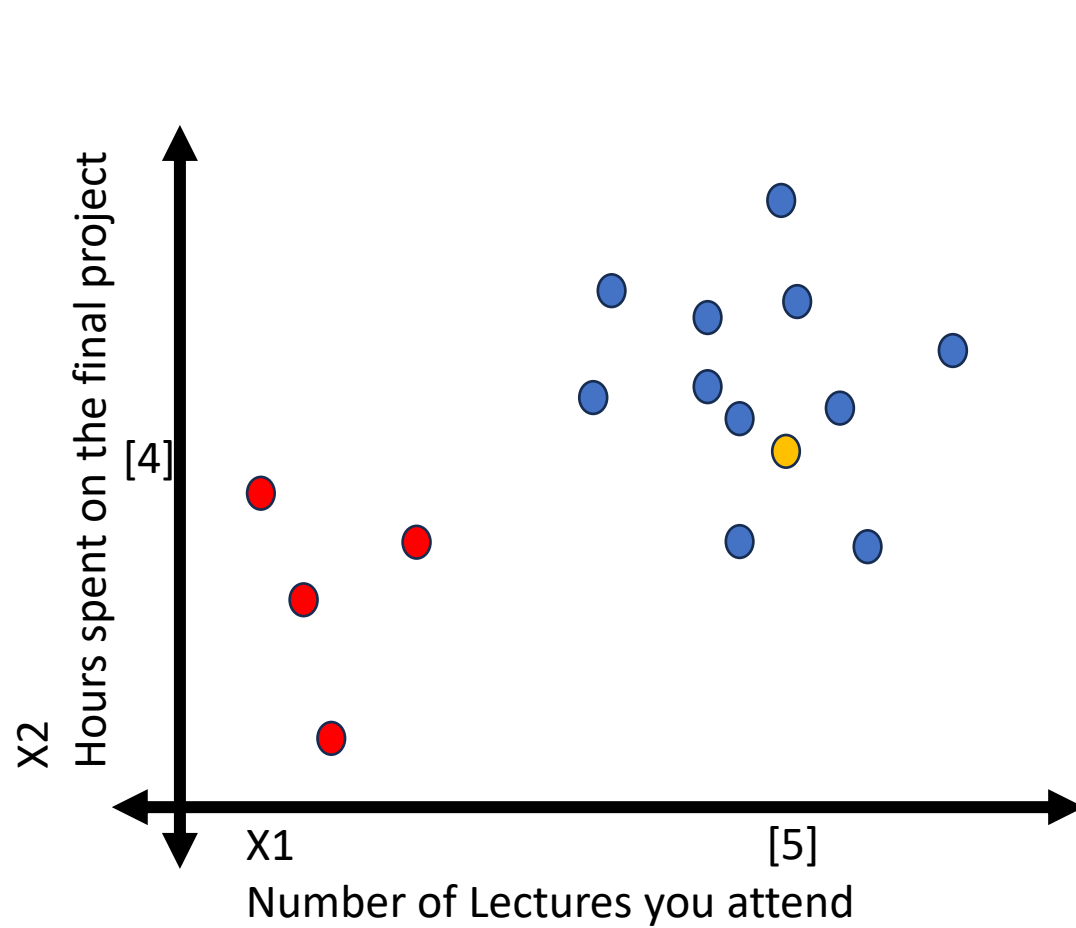
Toy Problem – Motivating Example

Two Feature Model:

X_1 = Number of lectures you attend

X_2 = Hours spent on the final group Project

WILL I PASS THIS CLASS?



Not trained properly! ← Why NN is so wrong? ← Predicted: $\hat{y} = 0.2$
 Actual: $y = 1$

$$\frac{1}{n} \sum_{i=1}^n \mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{W})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$

$f(x)$		y
0.1	✗	1
0.8	✗	0
0.6	✓	1
⋮		⋮

Loss function

Loss Optimization — Gradient Descent

1. Initialize weights randomly: $W \sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
 - (a) Compute gradient, $\frac{\partial J(W)}{\partial W}$
 - (b) Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
3. Return weights

Example in Python from scratch: NEURAL NETWORK MODEL

